

Computer Vision

(for Autonomous Driving)



Raoul de Charette



Texture segmentation


Data representation

- Find the richest most compact representation
 - Texture, Chromacity, Motion, Frequency, Entropy, etc.
- Solution is data-dependent



- Usually dealing with low saliency data

Single point model



Single point model




RGB [50 71 28]
HSV [90 154 71]



- Naïve solution: pixel-wise Euclidean distance (obviously fails)

Single point model



Single point model



RGB [50 71 28]
HSV [90 154 71]




- Naïve solution: pixel-wise Euclidean distance (obviously fails)



1. - RGBEuclideanDistance

Single point model



Single point model



RGB [50 71 28]
HSV [90 154 71]



Grass patch

- Naïve solution: pixel-wise Euclidean distance (obviously fails)




1. - RGBEuclideanDistance



1. - HSEuclideanDistance

Single point model



Single point model



RGB [50 71 28]
HSV [90 154 71]



Grass patch

- Naïve solution: pixel-wise Euclidean distance (obviously fails)



1. - RGBEuclideanDistance




1. - HSEuclideanDistance



1. - HSEuclideanDistance


Gaussian model



- Models the texture as a Multivariate Normal (3D Gaussian)
 - Advantage: models the signal variance
- **Probability Density Function (PDF)** allows computation of pixel-wise probability

$$\text{PDF: } f(\mathbf{x}) = \frac{1}{\sqrt{(2\pi)^k |\Sigma|}} \exp\left(-\frac{1}{2}(\mathbf{x} - \mu)^T \Sigma^{-1}(\mathbf{x} - \mu)\right)$$

Gaussian model



- Models the texture as a Multivariate Normal (3D Gaussian)
 - Advantage: models the signal variance
- **Probability Density Function (PDF)** allows computation of pixel-wise probability

$$\text{PDF: } f(\mathbf{x}) = \frac{1}{\sqrt{(2\pi)^k |\Sigma|}} \exp\left(-\frac{1}{2}(\mathbf{x} - \mu)^T \Sigma^{-1}(\mathbf{x} - \mu)\right)$$




PDF of RGB Normal



PDF of HSV Normal

Gaussian model




- Models the texture as a Multivariate Normal (3D Gaussian)
 - Advantage: models the signal variance
- **Probability Density Function (PDF)** allows computation of pixel-wise probability

$$\text{PDF: } f(\mathbf{x}) = \frac{1}{\sqrt{(2\pi)^k |\Sigma|}} \exp\left(-\frac{1}{2}(\mathbf{x} - \mu)^T \Sigma^{-1}(\mathbf{x} - \mu)\right)$$

- Why does it work so well ?




PDF of RGB Normal



PDF of HSV Normal

Gaussian model



- Models the texture as a Multivariate Normal (3D Gaussian)
 - Advantage: models the signal variance
- Probability Density Function (PDF) allows computation of pixel-wise probability

$$\text{PDF: } f(\mathbf{x}) = \frac{1}{\sqrt{(2\pi)^k |\Sigma|}} \exp\left(-\frac{1}{2}(\mathbf{x} - \mu)^T \Sigma^{-1} (\mathbf{x} - \mu)\right)$$

- Why does it work so well ?
- How will it perform to find these:



PDF of RGB Normal



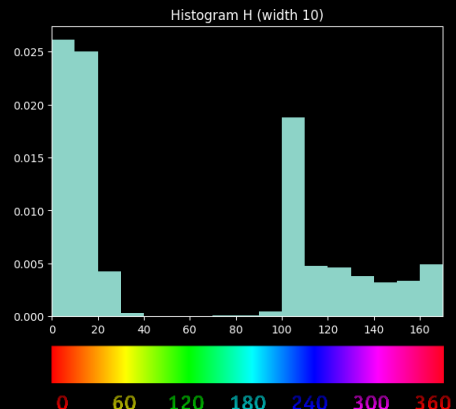
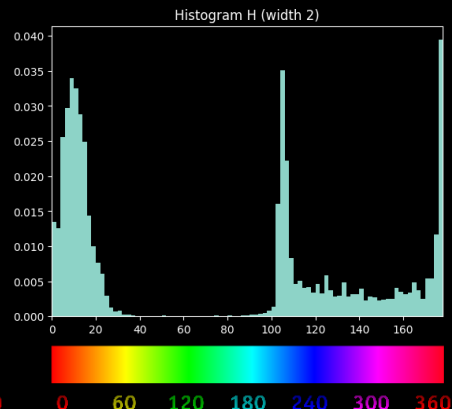
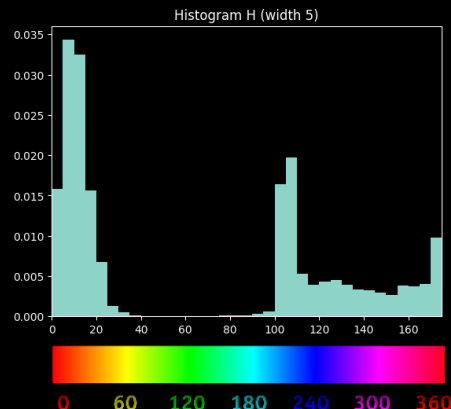
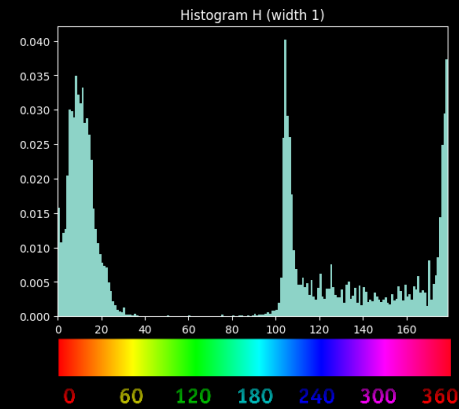

PDF of HSV Normal

Histograms

- Building histograms
 - Choose the right dimension-space
 - Usually no more than 2D for histograms
- Let's consider Histogram $\mathcal{H}(\dots)$ with h bins
- Bin stores number of occurrences

$$\mathcal{H}(\mathbf{x}) = |\{p \mid \forall p = \mathbf{x}, p \in I\}|$$

- Size matters for histograms.



- Back projection: likelihood of pixels to image histogram \mathcal{H}


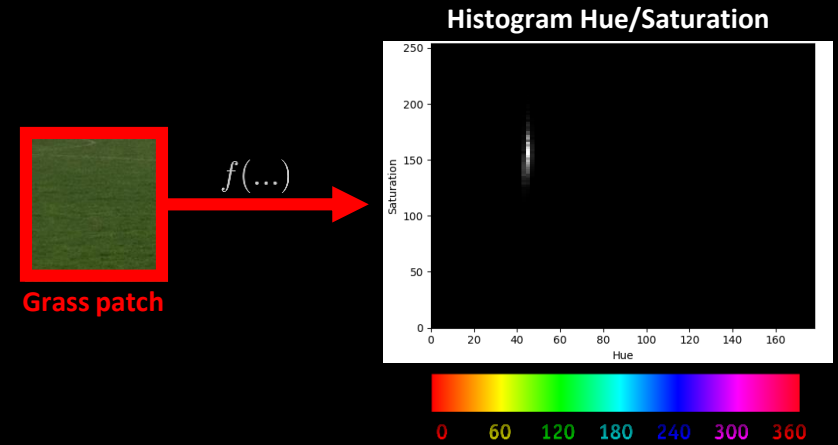
- Assume a normalized histogram: $\sum_{\mathbf{x}} \mathcal{H}(\mathbf{x}) = 1$

The normalized bins can be considered likelihood

Hence likelihood of pixel p to belong to histogram is:

$$\mathcal{H}(\underbrace{f(I(p))}_{\text{Mapped value}})$$

- Referred as: probability histogram



Back projection of grass Histogram

A words on histograms

- Bins can be asymmetrical => generally a bad idea

- Optimal bin size (h)

- Square root

$$n_h = \sqrt{n}$$

Naïve estimator

- Scott rule




$$h = \sigma \sqrt[3]{\frac{24 * \sqrt{\pi}}{n}}$$

Usually considered for large datasets. Not robust to outliers

- Freedman Diaconis Estimator (FD estimator)

$$h = 2 \frac{IQR}{n^{1/3}}$$

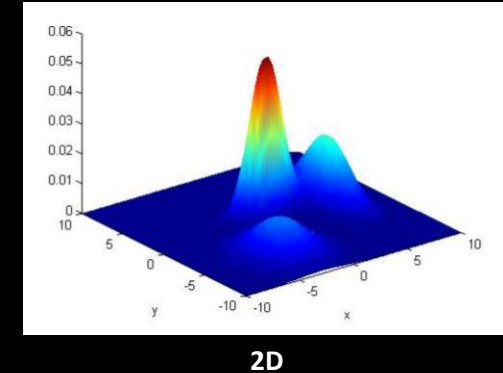

Using Interquartile Range (IQR). Optimal for large datasets. Robust to outliers.



Gaussian Mixture Models (GMM)

- GMM are really useful to represent mix of signals

$$\text{GMM} = \sum_i w_i \mathcal{N}(\mu_i, \sigma_i)$$




- Where the set of data X is supposed to be a partial image of the GMM
 - Models more complex signal
- Problem how to estimate the set of parameters $\{w_i, \mu_i, \sigma_i\}$?
 - Expectation Maximization, k-Means

Expectation-Maximization (Arthur P. Dempster)

$\{w_i, \mu_i, \sigma_i\}$ Gaussian Mixture Models (inc. noise)

- Estimates parameters of a statistical models partially observed
- How does it work ?
 - Tries to map a data to N Gaussians through maximization of likelihood
 - Randomly initialize a set of parameters (random Gaussians)
 - **E-Step**: For each point compute the probability to be generated by Gaussians
 - **M-Step**: Update parameters to maximize the (log) likelihood of the data
- Iterate E and M until convergence




(Mohand Saïd Allili, Tutorial, 2010)

Expectation-Maximization (Arthur P. Dempster)

$\{w_i, \mu_i, \sigma_i\}$ Gaussian Mixture Models (inc. noise)

- Estimates parameters of a statistical models partially observed
- How does it work ?
 - Tries to map a data to N Gaussians through maximization of likelihood
 - Randomly initialize a set of parameters (random Gaussians)
 - **E-Step**: For each point compute the probability to be generated by Gaussians
 - **M-Step**: Update parameters to maximize the (log) likelihood of the data
- Iterate E and M until convergence




(Mohand Saïd Allili, Tutorial, 2010)


Expectation-Maximization (Arthur P. Dempster)

$\{w_i, \mu_i, \sigma_i\}$ Gaussian Mixture Models (inc. noise)

- Estimates parameters of a statistical models partially observed
- How does it work ?
 - Tries to map a data to N Gaussians through maximization of likelihood
 - Randomly initialize a set of parameters (random Gaussians)
 - **E-Step**: For each point compute the probability to be generated by Gaussians
 - **M-Step**: Update parameters to maximize the (log) likelihood of the data
- Iterate E and M until convergence




(Mohand Saïd Allili, Tutorial, 2010)



Expectation Maximization (EM)

Which Gaussian does it belong ?



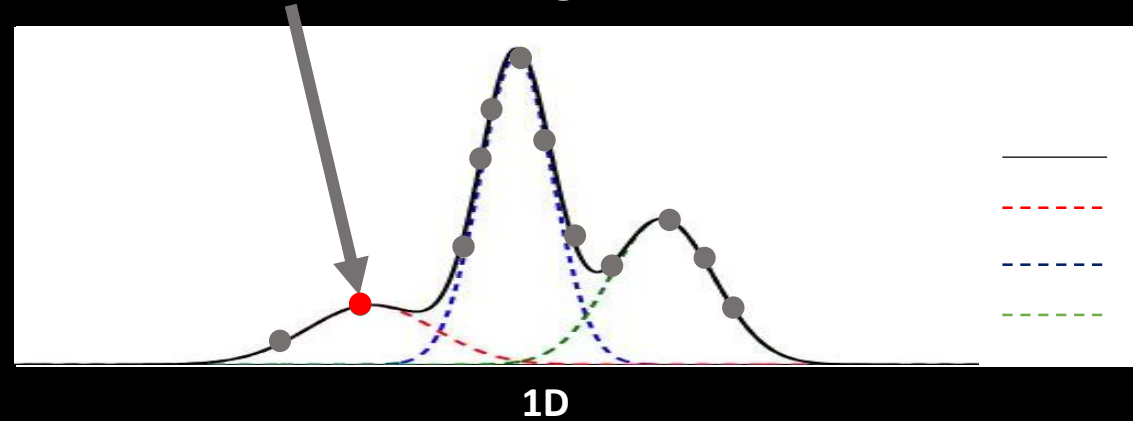
Suppose all parameters estimated $\{w_i, \mu_i, \sigma_i\}$

Assigning a label ("Gaussian") to each data points is straightforward

- The probability that a point belong to a Gaussian \mathcal{N} is: $p(x|\mathcal{N}) = \mathcal{N}_{PDF}(x)$
- Suppose n Gaussians, the class is the Gaussian with highest probability

$$\text{label}(x) = \operatorname{argmax}_i p(x|\mathcal{N}_i)$$

Which Gaussian does it belong ?




Suppose all parameters estimated $\{w_i, \mu_i, \sigma_i\}$

Assigning a label (“Gaussian”) to each data points is straightforward

- The probability that a point belong to a Gaussian \mathcal{N} is: $p(x|\mathcal{N}) = \mathcal{N}_{PDF}(x)$
- Suppose n Gaussians, the class is the Gaussian with highest probability

$$\text{label}(x) = \operatorname{argmax}_i p(x|\mathcal{N}_i)$$

Which Gaussian does it belong ?




Suppose all parameters estimated $\{w_i, \mu_i, \sigma_i\}$

Assigning a label ("Gaussian") to each data points is straightforward

- The probability that a point belong to a Gaussian \mathcal{N} is: $p(x|\mathcal{N}) = \mathcal{N}_{PDF}(x)$
- Suppose n Gaussians, the class is the Gaussian with highest probability

$$\text{label}(x) = \operatorname{argmax}_i p(x|\mathcal{N}_i)$$

Which Gaussian does it belong ?




Suppose all parameters estimated $\{w_i, \mu_i, \sigma_i\}$

Assigning a label ("Gaussian") to each data points is straightforward

- The probability that a point belong to a Gaussian \mathcal{N} is: $p(x|\mathcal{N}) = \mathcal{N}_{PDF}(x)$
- Suppose n Gaussians, the class is the Gaussian with highest probability

$$\text{label}(x) = \operatorname{argmax}_i p(x|\mathcal{N}_i)$$

Which Gaussian does it belong ?





Suppose all parameters estimated $\{w_i, \mu_i, \sigma_i\}$

Assigning a label ("Gaussian") to each data points is straightforward



- The probability that a point belong to a Gaussian \mathcal{N} is: $p(x|\mathcal{N}) = \mathcal{N}_{PDF}(x)$
- Suppose n Gaussians, the class is the Gaussian with highest probability

$$\text{label}(\mathbf{x}) = \operatorname{argmax}_i p(\mathbf{x}|\mathcal{N}_i)$$



GMM with EM









GMM with EM



3









GMM with EM





3



5









GMM with EM





3





5



17









GMM with EM





3




5



17



128



Summarizing


- Point Model are very naïve
- Gaussian Model can work well (e.g. keying)
- Histograms works well for texture with a few modes
- Gaussian Mixture Models are very powerful with distinct modes

Summarizing

- Point Model are very naïve
- Gaussian Model can work well (e.g. keying)
- Histograms works well for texture with a few modes
- Gaussian Mixture Models are very powerful with distinct modes
- But... Can you see a problem with all these methods ?
 - They all think at a pixel level
 - Spatial information is lost


Summarizing

- Point Model are very naïve
- Gaussian Model can work well (e.g. keying)
- Histograms works well for texture with a few modes
- Gaussian Mixture Models are very powerful with distinct modes
- But... Can you see a problem with all these methods ?
 - They all think at a pixel level
 - Spatial information is lost




Summarizing

- Point Model are very naïve
- Gaussian Model can work well (e.g. keying)
- Histograms works well for texture with a few modes
- Gaussian Mixture Models are very powerful with distinct modes
- But... Can you see a problem with all these methods ?
 - They all think at a pixel level
 - Spatial information is lost




Summarizing

- Point Model are very naïve
- Gaussian Model can work well (e.g. keying)
- Histograms works well for texture with a few modes
- Gaussian Mixture Models are very powerful with distinct modes
- But... Can you see a problem with all these methods ?
 - They all think at a pixel level
 - Spatial information is lost



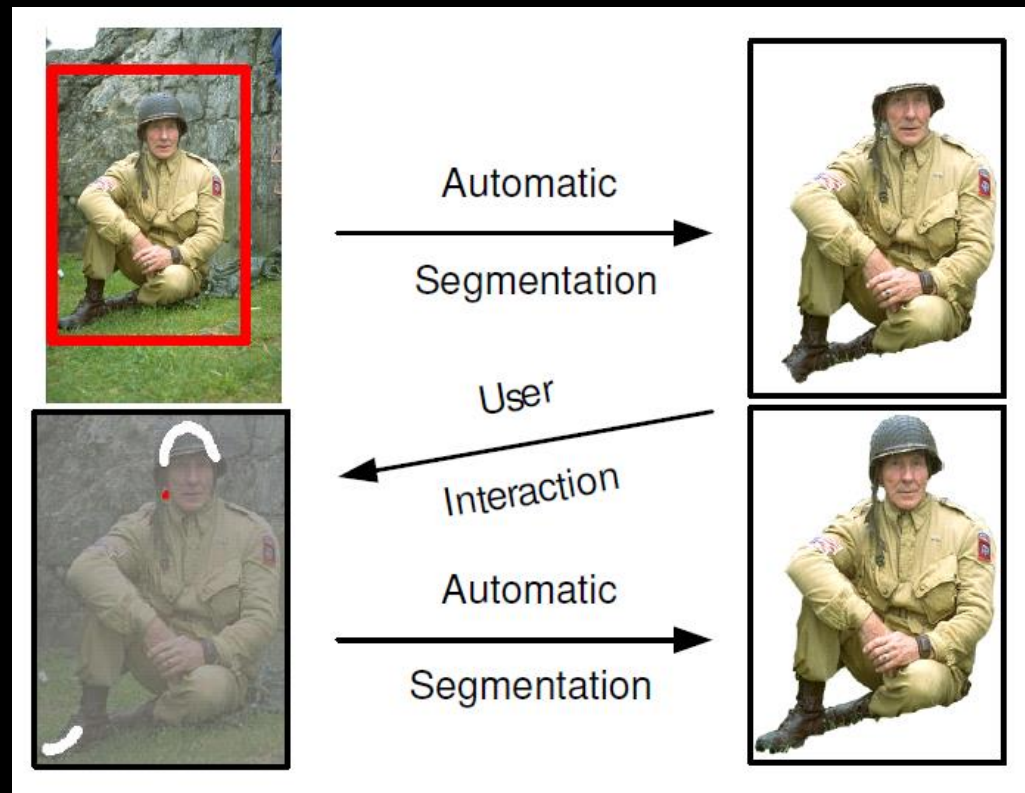
Summarizing

- Point Model are very naïve
- Gaussian Model can work well (e.g. keying)
- Histograms works well for texture with a few modes
- Gaussian Mixture Models are very powerful with distinct modes
- But... Can you see a problem with all these methods ?
 - They all think at a pixel level
 - Spatial information is lost




Interactive segmentation

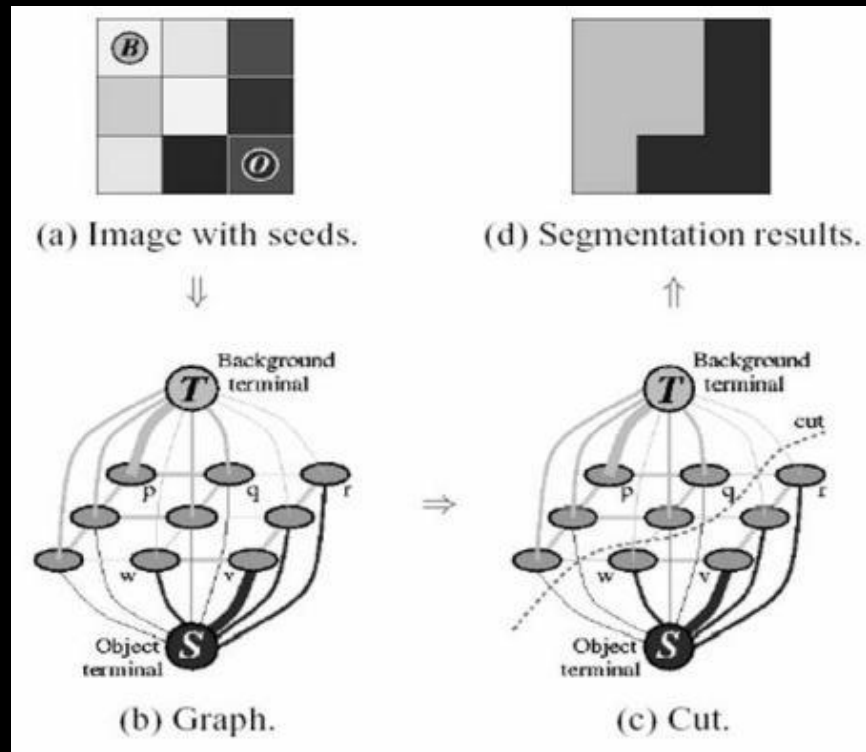
- What is it ?
- When to use interactive segmentation ?



- Region growing
 - Start from a seed
 - And iteratively grow if criteria is still valid
- Usual criteria:
 - Distance to seed $<$ threshold
 - Distance to neighbor $<$ threshold
 - Region (max-min) $<$ threshold
- Very sensitive to parameters
- Still used in 3D



- Grab cut



- Watershed






Image segmentation

Segment the **brick walls** (and else) with a histogram back projection

- Guidelines

Read image and display it

```
imBGR = cv2.imread("house.jpg")  
plt.imshow(imBGR[..., ::-1])
```

Select a sub region of the image to be the template and display it

```
patchBGR = imBGR[y1:y2, x1:x2]  
plt.imshow(patchBGR[..., ::-1])
```

Convert BGR patch to HSV. (use `cv2.cvtColor` and `cv2.COLOR_BGR2HSV`)

Build the histogram:

```
hist = cv2.calcHist([patchHSV], channels=[0], mask=None,  
                   histSize=[180], ranges=[0, 180])
```

and normalize it (divide by max value)

Compute the back projection of the histogram

```
cv2.calcBackProject([imHSV], channels=[0], hist=hist, ranges=[0, 180], scale=255)
```

- Once walls are segmented continue with: **grass, roof, pathway**

- Finished ?** 1) Build a real segmentation map where 0=void, 1=brick, 2=roof, 3=grass, 4=path. 2) Apply Gaussian model segmentation for sky.
- You're really good ?** 1) Code your own back projection, start using `calcHist`. 2) Use a FD estimator for hist (yes, you need to recode `calcHist`). Helps others 😊
- You're even better ?** Cluster the image with a GMM model (say: 20 Gaussians) and assign grass/roof/wall/pathway/sky label to each.

